

Case Builder — Six-Day Execution Packet

Founder desk reference · Mon 2026-08-03 → Sun 2026-08-09 · Cloud QA only · Production & waterwoods untouched

Page 1 — Problem, promise, reality, claims

Customer problem. Chinese-speaking customers and small broker offices still run accidents through phone calls, scattered WeChat messages, repeated document asks, unclear ownership, and weak case continuity.

North Star. Recognize known customer → reuse confirmed vehicle/policy → collect story & evidence → find real gaps → Request More → Broker confirms → one Active Case / Timeline / next action → measure timing and AI quality. Do not sell LangGraph, LangSmith, or MCP.

Current product reality

Area	Status
Stage 1 Claim path	Founder phone validated (frozen)
Stage 2 known-customer confirm	Founder phone validated (closed)
Timing metrics V1	Closed on QA (honest, small-n)
LangGraph accident-story assist	Code + tests + on QA · phone freeze pending
LangSmith Case Builder eval	Not started
AI accept/edit/reject rates	Not implemented yet
MCP Broker tools	Lab only · not this week's P0
Production release	Not authorized

What you can / cannot claim

Can say	Cannot say
Founder-demonstrable on Cloud QA	Production ready
Truthful V1 / supervised soft-pilot talk (with restrictions)	Real-customer validated
AI proposes; humans confirm	Autonomous claim filing AI
We measure timing honestly	Proven ROI / minutes saved
Eval & safety work this week	Legal/compliance certified

Page 2 — Six-day schedule & Founder actions

Day	Focus	Exit gate	Founder action
Mon 8/3	Freeze LangGraph PR A	Auto green + phone result	≤5 min phone QA
Tue 8/4	LangSmith traces + golden eval	Offline eval gate green	None (judge optional later)
Wed 8/5	AI feedback + latency metrics	Events export; no ROI claims	Read 1-page metrics note
Thu 8/6	Pilot safety baseline	Audit + fallback checklist	Sign audit / consent intent (~15m)
Fri 8/7	One-click demo + portfolio draft	Chen path works	Run one-click; skim pitches
Sat 8/8	Portfolio polish	Evidence index complete	Approve public wording
Sun 8/9	Final gate	Truthful label + demo clip	One E2E walkthrough; pick label

P0 before any MCP: LangGraph freeze → LangSmith eval → AI metrics → safety → demo/portfolio → Sunday gate.
MCP = P1 only after P0s. Optional LLM judge / polish / cost cleanup = P2.

Daily habit: Capture one evidence artifact every day (do not wait for Saturday).

Page 3 — Workflow & how modules connect

Complete workflow: Customer entry → known-customer context → accident story → AI proposal → customer confirmation → completeness check → Broker review → Request More (if needed) → customer supplement → office acceptance → Timeline / metrics

Modules (customer value)

Module	Role	Customer / office value
Deterministic Claim workflow	Backbone	Repeatable office process
Timeline & projections	One truth	Customer & Broker see the same case
Known-customer confirm	Prefill gate	Fewer unnecessary uploads
LangGraph	Story organizer	Faster structured Must Haves
LangSmith	Flight recorder + quizzes	Catch bad AI before brokers see it
Metrics	Honest clocks	Learn where cases stall
HITL feedback	Accept / edit / reject	AI improves without guessing
Safety & audit	Boundaries	Soft pilot talk without reckless risk
Model fallback	Safety net	Form still works if AI fails
MCP (later)	Delivery tools	Second-office leverage — not the sale

Control law: AI never changes Claim lifecycle. Unconfirmed proposals are not facts.

Page 4 — Stories & Sunday checklist

Three-minute customer story (for Chen / office owner)

After an accident, your customer opens one task—not a chat scavenger hunt. If they are a known customer, they confirm the car and policy instead of re-uploading everything. They tell the story once; the system organizes it and asks only the missing pieces. Your office sees one case, one timeline, and one next action—Request More when needed, then accept when materials are ready. We are not filing with the carrier or deciding liability. We are making the office-ready case faster and clearer.

Three-minute FDE interview story

We shipped a deterministic claim workflow first and Founder-validated it on phone. Then known-customer confirm. Then honest timing metrics that never invent broker-open events. Only after that did we add a bounded LangGraph assistant that proposes structured accident facts—humans must confirm; the graph never mutates lifecycle; invalid model output falls back to rules. This week we add LangSmith evals, Accept/Edit/Reject metrics, and a pilot safety baseline. We sell fewer loops and clearer ownership—not a multi-agent platform.

Sunday release checklist

- ☐ Critical automated tests green
- ☐ One Founder E2E walkthrough on QA
- ☐ Fix only P0 blockers
- ☐ Short demo recorded
- ☐ QA cost-saving scale restored
- ☐ Evidence index updated
- ☐ Label: DEMO READY / PILOT READY WITH RESTRICTIONS / NOT PILOT READY
- ☐ Production & waterwoods still untouched

Note: PDF uses Latin-safe fonts (no CJK on this host). For full Chinese print, open the companion HTML or markdown. No secrets or customer PII.